A04 Analyzing "Arrival" Through the Lens of NLP

Jiri Musil

Department of Science, Technology, Engineering & Math,

Houston City College

ITAI 2376 Deep Learning Artificial Intelligence

Professor Patricia McManus

February 12th, 2026

Introduction

Arrival (2016) is a science fiction film in which the central “alien invasion” is really a communication problem. Twelve new spacecraft land around the world, and the risk stems from uncertainty about the visitors' intentions, with minor misinterpretations swiftly escalating into political conflicts. That tension serves as a valuable reflection for Natural Language Processing (NLP). In actual text and speech, meaning is seldom a straightforward one-to-one correlation. Words have multiple senses, speakers imply more than they say, and culture shapes what is considered polite, threatening, or humorous. This report highlights key scenes from Arrival to illustrate how communication breakdowns in the movie relate to real-world NLP challenges and how the team’s methods align with rule-based, statistical, and modern neural approaches.

(*What We’ve Got Here: “Arrival” | Los Angeles Review of Books*, 2016).

Body

The film illustrates several NLP challenges that also appear in real-world language technology.

1. Low-resource learning and grounding

Banks is effectively working in a low-resource setting as there is no dictionary, no grammar, and no large body of examples to draw on. The initial task is not complex translation but rather gathering grounded data to establish shared reference points visible to both parties. Her choice to begin with concrete labels and controlled interactions mirrors a core lesson in NLP, where models (and people) learn faster when language is tied to observable context rather than just symbols on a page. It also reflects a wider critique in NLP research, where pattern forms are not equivalent to meaning unless they are linked to real contexts and intentions.

2. Segmentation, representation, and structure

Even after communication begins, the logograms pose a familiar technical question: what are the basic units? While human readers typically depend on spaces and word boundaries, the heptapods create a single circular mark that encodes an entire thought. Before translation can occur, Banks requires processes similar to tokenization and parsing, which involve recognizing recurring sub-patterns and determining what is considered stable structure. The film's designers deliberately incorporated recurring elements into the logograms and utilized computational tools to analyze patterns, treating the writing as analyzable data.

(Rhodes, 2016).

3. Ambiguity and word sense under pressure

The “offer weapon” crisis is a clear parallel to word-sense ambiguity, a form with multiple plausible meanings. In everyday NLP applications such as machine translation and question answering, selecting the wrong sense can completely change a sentence's meaning. Arrival adds an important human layer, namely ambiguity, interpreted through fear, incentives, and prior beliefs. When decision-makers are set up for conflict, their most aggressive instincts can seem obvious, even if the language evidence is weak. This is why real-world NLP systems need careful error analysis and clear communication of uncertainty, especially when outputs feed into high-stakes decisions.

4. Pragmatics, idioms, sarcasm, and cultural variation

Most of the film’s challenge is more practical than just meaning as the team is always trying to figure out purpose and intent, not just decode. Pragmatics, the way meaning relies on context, shared assumptions, and unspoken implications is a recognized weakness for language technology. Arrival also hints at how culture complicates interpretation. The heptapods are nicknamed “Abbott” and “Costello,” a reference that relies on shared cultural knowledge and a comedy routine built on name misunderstandings. In real text, the same problem appears with idioms and sarcasm, where the literal meaning is not the real message. Idioms are usually not made up of parts that fit together easily, and sarcasm often needs social context to be understood. The film shows “domain shift" literally, showing how countries interpret translations differently, causing misunderstandings to grow when information sharing breaks down. Arrival dramatizes the cost of assuming that meaning is universal when the background assumptions are not shared.

The team’s workflow also mirrors several major NLP approaches.

1. Rule-based: controlled protocols and interpretable mappings

Banks’ early work looks “rule-based” in spirit. She limits variables, builds a small inventory of reliable mappings, and refuses to leap into complex questions without groundwork. This reflects traditional symbolic NLP, where meaning is captured through explicit rules and structured formats. This is beneficial because the reasoning process can be examined and challenged. The trade-off is speed and brittleness because hand-built mappings take time and can fail when the language behaves unexpectedly.

2. Statistical: pattern discovery and alignment from examples

As the logogram “corpus” grows, the work naturally shifts toward statistical thinking, including cataloging examples, comparing them, and revising hypotheses based on recurring patterns. Arrival also warns that statistics alone are not enough in a new domain. With limited or skewed data, you can become confident in an incorrect generalization, which is one reason domain adaptation remains a major practical challenge.

(Rhodes, 2016).

3. Neural NLP and hybrid toolkits

Modern neural methods could potentially help in pattern recognition in the raw signals such as audio, video, and the developing parallel data linking logograms and English explanations. However, the film emphasizes a crucial limitation , where recognizing patterns in form does not equate to grasping meaning and intent. The most convincing "real-world" solution is a hybrid with effective data capture and annotation, models capable of proposing hypotheses, and humans testing those hypotheses through interaction. In essence, the tool goes beyond being merely a model, it is a process that incorporates verification, uncertainty tracking, and shared context.

Innovative perspective

The movie made me think about possible approaches to dealing with its challenges and also prompted reflections as a human resources professional and on my experience with communication. One approach could be as follows. A practical NLP system would treat translation as an interactive, risk-aware protocol rather than a single best-guess output. For high-stakes phrases (e.g., words that could imply violence), the system should return the top 3 interpretations, each with a confidence score and a brief justification based on context. If confidence is low or the stakes are high, the system should refuse to finalize and instead generate a follow-up question to reduce uncertainty (an active-learning loop). This approach mirrors Banks' strategy of establishing shared referents, testing hypotheses, and expanding the "vocabulary" only when the evidence is stable. Finally, because Arrival shows that different groups can interpret the same message differently, the system should include cross-cultural evaluation and monitoring to check whether outputs shift in meaning across audiences and domains and to adapt accordingly.

Conclusion

Arrival's most useful contribution to thinking about NLP is how it treats misunderstanding as normal and communication as collaborative. Banks succeeds by creating common ground and by correcting course when early assumptions prove wrong. The movie also makes a practical point about deployment, noting that translation outputs do not stay on a whiteboard. They shape decisions, emotions, and politics, and the consequences of errors depend on the setting. This highlights a challenge in NLP evaluation, where the focus is often on average accuracy, but the potential for catastrophic worst-case mistakes is overlooked.

References:

1. Brower, J. (2016, December 12). What We've Got Here: "Arrival". Los Angeles Review of Books. https://lareviewofbooks.org/article/weve-got-arrival/
2. Coon, J. (2019). The linguistics of Arrival: Heptapods, field linguistics, and Universal Grammar. Manuscript. https://jessica.lingspace.org/wp-content/uploads/2019/04/Coon%E2%80%93Arrival.pdf
3. Jurafsky, D., & Martin, J. H. (2026). Speech and language processing: An introduction to natural language processing, computational linguistics, and speech recognition with language models (3rd ed.). Online manuscript released January 6, 2026. Stanford University. https://web.stanford.edu/~jurafsky/slp3/
4. O'Callaghan, J. (2016, November 12). The science of Arrival: what the film got right (and wrong). WIRED. https://www.wired.com/story/arrival-science-fact-fiction/
5. Ramponi, A., & Plank, B. (2020). Neural unsupervised domain adaptation in NLP—A survey. Proceedings of the 28th International Conference on Computational Linguistics, 6838–6855. https://aclanthology.org/2020.coling-main.603/
6. Rhodes, M. (2016, November 16). How Arrival's designers crafted a mesmerizing alien alphabet. WIRED. https://www.wired.com/2016/11/arrivals-designers-crafted-mesmerizing-alien-alphabet/
7. Salton, G. (2017). Representations of idioms for natural language processing: Idiom type and token identification, language modelling and neural machine translation (Doctoral dissertation, Technological University Dublin). https://arrow.tudublin.ie/sciendoc/201/

8. Zhou, M., Duan, N., Liu, S., & Shum, H. (2020). Progress in neural NLP: modeling, learning, and reasoning. *Engineering*, *6*(3), 275–290. https://doi.org/10.1016/j.eng.2019.12.014